

Who Discovers the Price?

A real cross-venue price-discovery lead — and why a market-maker can't trade it. Tick-level evidence from the 2026 World Cup.

Prabhat M · @PrabhatM27 · Math & Statistics, Purdue · github.com/Thesavagecoder7784/xResidual · 23 June 2026 · working note


*The headline of this note is a non-result, and that is the point. Using tick-level order-book data from 24 World Cup matches on Kalshi and Polymarket, I find a genuine cross-venue price-discovery lead — on a goal, Polymarket's mid moves ~400 ms first and carries a ~78% Gonzalo–Granger information share, leading 20 of 21 cointegrated matches (computed on mid-quotes, which sidesteps the ~59% trade-direction problem documented for Polymarket). Conditioned on the news the lead is sharper still: ~86% of price discovery in goal windows versus ~53% in calm play. **And then it is not tradeable.** At the goal the book vanishes — spread ~8× wider, best-price depth under 1% of normal — so the ~11-cent stale-quote gap a reactor would lift nets +10c on paper but 0% after gating on the depth actually resting in the book. That collapse-and-refill is the market-maker pulling quotes against toxic, information-motivated flow: **adverse selection observed in real time**, and the reason the lead is a property of price discovery, not a free trade. Forecasts for the wider project were pre-registered before kickoff and graded publicly on 19 July — verifiable without trusting the author.*

Summary.

- **Claim:** Polymarket is the price-discovery leader of the Kalshi–Polymarket pair on the same World Cup outcomes.
- **Independent methods agree:** an event-study reaction lead, Hasbrouck / Gonzalo–Granger information shares, an order-flow mechanism check, and a cross-sectional level check against the sharp line.
- **Numbers:** lead-lag **103 vs 45** events to Polymarket (**62%**, median **+400ms**) across 34 matches; Gonzalo–Granger share **~78%**, leading **20/21** cointegrated matches; OFI→price impact **t≈58** within-venue; mean gap to the sharp line **0.18pp** (Poly) vs **0.26pp** (Kalshi).
- **Anatomy of a goal:** discovery concentrates at the news (Polymarket **~86%** in goal windows vs **~53%** calm); the book withdraws hardest where it leads (spread **~8×**, depth **<1%**); so the ~11c stale-quote gap is +10c on paper but **0%** harvestable once gated on the vanished depth.
- **Honest caveats:** not tradeable after cost (sub-second + the book withdraws); one tournament, one sport; event-clustering means the formal significance unit is the match, not the event.
- **The maker's view:** the collapse-and-refill is the quoter's defence against toxic flow — a maker who held stale quotes through the goal would be the one adversely selected. The unharvestable gap is exactly the loss the withdrawal avoids.
- **Verifiable, not trust-me:** the project's forecasts were pre-registered before kickoff and are graded publicly on 19 July, hits and misses — a track record that can't be edited after the fact.
- **Why it matters:** it answers a question a 2026 Polymarket microstructure study explicitly left open (on live in-play sports, not the election markets prior work used), and it tells anyone aggregating prediction-market signals which venue to trust first — while showing, with the costs in, that the lead is information not alpha.

A goal is a tax on whoever's still *quoting*.

When a goal hits, fair value jumps $\sim 11c$ in a few hundred milliseconds. A market-maker who left a quote **resting** gets lifted at the old price and is instantly down $\sim 11c$ per contract — picked off by exactly the flow that saw the goal. So makers **pull**: spread $8\times$ wider, depth to near zero, refilling in $\sim 3s$. The liquidity withdrawal isn't a glitch, it's the **adverse-selection defence** — and the same reason the price-discovery lead can't be traded: there's nothing left to lift.



illustrative single goal on a 50-50 contract; magnitudes are the pooled medians — stale-quote gap $\sim 11.4c$, spread $\sim 8\times$, depth $\sim 1\%$, refill $\sim 3s$ (~ 378 goal shocks, ~ 28 matches) @PrabhatM27 / xResidual

Figure 1. The anatomy of a goal, from the maker's side. Fair value jumps $\sim 11c$ in a few hundred milliseconds. A quote left resting in the book is lifted at the stale price and is down $\sim 11c$ per contract — so makers pull: spread $\sim 8\times$ wider, depth to near zero, refilling in $\sim 3s$. That withdrawal is the adverse-selection defence, and the reason the price-discovery lead a reactor sees is information, not a tradeable edge. (Illustrative single goal on a 50-50 contract; magnitudes are the pooled medians across ~ 370 goal shocks.)

1. The open question

Kalshi (US-regulated) and Polymarket (crypto-native) quote the *same* World Cup outcomes through very different microstructure — different fee and tick regimes, user bases, and settlement. That makes the pair a natural laboratory for cross-venue price discovery: when new information arrives, which venue moves first, and which merely follows?

A 2026 study of the Polymarket order book [1] maps that venue's microstructure in detail but states plainly that a cross-venue comparison against Kalshi — the price-discovery question — is left for future work. This note answers it for a clean, repeated, exogenous information event: goals in the 2026 World Cup. I form no strong prior on direction and let the data decide.

2. Data

I capture both venues' full level-2 order books and every trade over WebSocket, at tick resolution, from an always-on collector running 24/7 through the tournament. For each match I pair contracts that reference the identical outcome (one Kalshi ticker and one Polymarket token per team and the draw), de-vig each venue's quotes to probabilities, and align the two mid-price series on a common clock. The event study and OFI run on **34 matches**, and the information-share estimates (below) on the **21 cointegrated matches** (34 contracts) whose cross-venue spread passes the ADF guard (all as of 23 June).

3. Methods

I measure the same phenomenon three ways, at increasing rigor, so the conclusion does not rest on any single technique.

3a. Event-study reaction lead

I detect repricing shocks (goals) on each venue's mid series and cross-correlate the two around each event to estimate which venue's mid moves first, to the millisecond. Intuitive, but timing in a sub-second window is noisy.

3b. Information shares (the core)

The two de-vigged mids for one outcome are cointegrated with the natural vector $(1, -1)$: arbitrage keeps their difference stationary. I test that with an Augmented Dickey–Fuller test and *discard any contract that fails* (on a non-cointegrated pair the shares are a spurious-regression artifact). On the survivors I fit a bivariate vector error-correction model and report the **Hasbrouck** [3] information-share bounds and the ordering-free **Gonzalo–Granger** [4] permanent-component share. The intuition: the venue that adjusts *less* to the spread (the smaller error-correction coefficient) is the one others converge toward — the price-discovery leader.

3c. Mechanism

I reconstruct order-flow imbalance (OFI) from top-of-book changes [5] and test (i) whether OFI moves each venue's own price — the canonical within-venue result — and (ii) whether one venue's flow predicts the *other's* next move. I also recompute the lead on the imbalance-weighted micropurchase [6] rather than the mid.

4. Results

The winner market is efficient. As a control, the de-vigged outright-winner prices agree across venues (and with a calibrated model) to within ~ 0.4 pp. There is no cross-venue mispricing in the liquid market; the effect below is about *speed of discovery*, not a level disagreement.

And where they do differ, Polymarket sits closer to the sharp line. Benchmarking each venue's de-vigged title price against the Betfair exchange (the standard sharp anchor) across 48 teams, Polymarket's mean absolute gap to the sharp line is **0.18 pp** versus Kalshi's **0.26 pp**, and Polymarket is the closer of the two on 22 teams to Kalshi's 11. So the venue that updates *first* is also the one sitting *closest* to the sharpest available consensus: here, leadership and level accuracy coincide — a mild counterpoint to the “leadership need not mean efficiency” caution raised for 2024 US-election markets.

MEASURE	SAMPLE	RESULT
Reaction lead (event study)	34 matches, 165 events	Polymarket first in 103 , Kalshi in 45 (62%), median +400 ms (IQR [-200, +800])
Gonzalo–Granger share	21 cointeg. matches	Polymarket ~78% , leads in 20 / 21
Hasbrouck share (mid of bounds)	17 contracts	Polymarket ~63% (band 59–71%)
Distance to sharp line (Betfair)	48 teams	Polymarket 0.18pp vs Kalshi 0.26pp ; Poly closer on 22 / 11
OFI → own-venue price impact	17 matches	strong, both venues (Kalshi $t \approx 29$, Poly $t \approx 35$)
OFI cross-venue lead	17 matches	null — no clean asymmetry either direction

All three methods point the same way: Polymarket leads. The mechanism check sharpens the interpretation — order flow clearly drives price *within* each venue (textbook [5]), but there is *no* cross-venue flow lead. The lead therefore lives in the *price*, not in observable flow: Polymarket's mid simply updates to new information first.

5. The anatomy of a goal

The results above are pooled over each match. Conditioning on the actual information event — a goal — turns the single ~78% into a dynamic picture, and answers the question any price leader should be asked: is the lead a tradable edge?

MEASURE (PER GOAL)	SAMPLE	RESULT
Information share, event-conditional	~20 matches	Polymarket GG ~86% in goal windows vs ~53% in calm play (Hasbrouck ~82% vs ~56%)
Liquidity at the goal (top of book)	~370 shocks	spread ~8× (Poly) / ~2× (Kalshi); best-price depth falls to <1% of normal; refill ~3–4s
Harvestability of the lead	117 goals	gross stale gap ~11.4c, net +10.2c on paper after spread — but 0% harvestable once gated on available depth

Discovery concentrates at the news. Away from goals the two venues are near-coequal (~53/47); the instant a goal lands, Polymarket carries ~86% of the permanent price move. The flagship lead is not a steady background hum but a *news-event* phenomenon — Polymarket leads precisely because it incorporates the goal first.

And the same venue withdraws liquidity the hardest. At the goal, makers pull their quotes: the spread blows out (~8× on Polymarket, ~2× on Kalshi) and best-price depth collapses to under 1% of normal, refilling only after a few seconds. Unlike a betting exchange — Betfair suspends and

cancels resting orders on a goal (Croxson & Reade 2014) — Kalshi and Polymarket never halt, so this is *informational* liquidity withdrawal, not a mechanical suspension.

So the lead is real but untradeable. A naive ledger makes the stale follower look like free money: ~11.4c of staleness, +10.2c net of the spread, on every goal. But that quote has vanished — depth at the goal is ~0.4% of normal — so once gated on what is actually resting in the book, 0% is harvestable. "The news is in the price before you can trade it" (Croxson & Reade 2014) holds here for a precise, measured reason: not slow pricing, but a book that empties in the same instant it reprices. (An earlier version of this ledger reported the +10.2c as harvestable; gating on depth corrected it, and that correction is the result.)

6. Robustness

Immune to the trade-direction problem. [1] shows that trade direction inferred from Polymarket's public feed agrees with on-chain ground truth only ~59% of the time, which corrupts any signed-trade measure (effective spread, Kyle's λ , trade-sign OFI). Every estimate here is computed on *mid-quote* moves and book changes, not signed trades, so it is unaffected; I verified our signed-flow code is used by none of these results.

No spurious shares. The ADF gate drops non-cointegrated pairs, so a Hasbrouck/GG number is only reported where the (1, -1) spread is genuinely stationary.

Significance, stated honestly. Goal events cluster within matches and are not independent, so the naive binomial ($p \approx 0.0004$) and bin-level OFI t-statistics overstate significance. The unit that carries the claim is the *match*: 20 of 21 cointegrated matches show a Polymarket lead, with a ~78% central share. A formal write-up would replace the naive p-value with a match-clustered bootstrap; the qualitative result is not in doubt.

7. What this is not

Not a trade. The lead is a few hundred milliseconds. Documented cross-venue arbitrage exists (~\$40M, 2024–25) but is seconds-lived and fee-bounded; a 400 ms head start vanishes inside the bid–ask once costs are crossed. This is price-discovery *leadership*, a property of market structure — not free money.

Scope. One tournament, one sport, a ~three-week window, 20 matches. Polymarket's documented wash trading (~25% of historical volume) is a caveat on volume-based measures, less so on mids. The Hasbrouck band is wide where the two venues' innovations are correlated, which is why I lead with the ordering-free Gonzalo–Granger number.

8. Related work & contribution

The within-venue OFI→price relation reproduces Cont–Kukanov–Stoikov [5]; the microprice's edge over the mid follows Stoikov [6]; the information-share machinery is Hasbrouck [3] and Gonzalo–Granger [4]; an independent 2025 study finds a Polymarket discovery lead in other markets [2],

which this corroborates — but where [2] reads that lead as “economically meaningful arbitrage,” the depth-gated ledger here qualifies it: on a clean repeated-event stream the lead is real, yet the book withdraws in the same instant it reprices, so the stale-quote gap is ~0% harvestable and the arbitrage is *not* economically meaningful net of the cost of immediacy. The contribution is, to my knowledge, the first **cross-venue Kalshi–Polymarket information-share estimate at tick resolution on a clean, repeated event stream** (goals), robust to the trade-direction problem [1], paired with a depth-gated harvestability test that separates a real price-discovery lead from a *tradable* one.

9. Conclusion

On identical World Cup outcomes, Polymarket discovers price first: it reacts ~400 ms sooner and carries ~78% of the permanent price movement, in 20 of 21 cointegrated matches measured. The practical implication is narrow but real — an aggregator of prediction-market signals should weight Polymarket's mid as the leading indicator and Kalshi's as the lagging confirm — and the trading implication is honestly none, after costs. Everything is reproducible from the public repository; forecasts in the wider project are committed to a forward-only ledger, and the core estimation math is frozen, so none of this is fit after the fact.

References.

1. "The Anatomy of a Decentralized Prediction Market: Microstructure Evidence from the Polymarket Order Book," 2026, arXiv:2604.24366.
2. Ng, Peng, Tao, Zhou, "Price Discovery and Trading in Modern Prediction Markets," 2025 (SSRN 5331995).
3. Hasbrouck, J. (1995), "One Security, Many Markets: Determining the Contributions to Price Discovery," *Journal of Finance*.
4. Gonzalo, J. & Granger, C. (1995), "Estimation of Common Long-Memory Components in Cointegrated Systems," *JBES*.
5. Cont, R., Kukanov, A. & Stoikov, S. (2014), "The Price Impact of Order Book Events," *Journal of Financial Econometrics*.
6. Stoikov, S. (2017), "The Micro-Price: A High-Frequency Estimator of Future Prices," SSRN 2970694.

xResidual — a live study of how prediction markets price the 2026 World Cup. Methods, data pipeline, and the live findings dashboard: github.com/Thesavagecoder7784/xResidual